

Conway's Game of Life Implementation



University of Sri Jayewardenepura
Faculty of Computing

B.comp (Hons.) in Computer Science
CCS3032 – High Performance Computing

Requested by: Dr. Buddhitha Hettige

Name: M.A.O. Rodrigo
Reg. No: FC115278
Index No: FC221019
Group No: 02

Table of Contents

1. Introduction.....	2
2. Problem definition	2
3. Objectives	2
4. Experimental setup	3
4.1 Platform.....	3
4.2 Grid and simulation parameters	3
4.3 CPU implementation.....	3
4.4 GPU implementation	3
4.5 Benchmarking methodology	3
4.6 Correctness verification	4
5. CPU implementation	5
6. CUDA implementation	8
7. Comparison results	12
8. Conclusion	13

1. Introduction

Conway's Game of Life is a cellular automaton where every cell in a grid is either alive or dead and the grid evolves in discrete generations in accordance with a fixed rule based on each cell's neighbors. Computing large grids over many generations for this implementation is computationally expensive because every cell must be evaluated independently at every generation. Therefore, we can use High-Performance Computing techniques to implement this.

In here I have implemented Game of Life in 2 methods,

- Conventional sequential CPU implementation using 2D arrays
- Parallel GPU implementation using NVIDIA CUDA

Both implementations are benchmarked under identical conditions to evaluate how much performance is gained by the parallel nature of the problem.

2. Problem definition

Each cell in an $n \times n$ grid follows Conway's standard rules for each generation.

- A live cell with fewer than 2 or more than 3 live neighbors dies (Under and Overpopulation)
- A live cell with 2 or 3 live neighbors survives
- A dead cell with exactly 3 live neighbors becomes alive (Reproduction)

Computing one generation requires visiting every cell and counting its up-to-8 neighbors. It gives $O(n^2)$ work per generation. For 'g' generations it will be $O(n^2 \times g)$. In a single CPU core this is done sequentially. The problem is how much faster this can be done through mapping each cell to an independent GPU thread and to quantify and explain the resulting performance difference in a statistically reliable way.

3. Objectives

- Implement Game of Life on a 1024×1024 grid using standard 2D arrays in C++ which is executed sequentially on CPU.
- Implement Game of Life on a 1024×1024 grid using standard CUDA which maps cells to GPU threads for parallel execution.
- Ensure both implementations use identical initial conditions such as fixed random seed, same density and identical rules so the results are directly comparable.
- Benchmark both implementations with repeated timed trials with an untimed warm up pass beforehand and report the mean and range. This reduces the influence of one-off effects such as cold caches or first-run allocation overhead that you may encounter in a single measurement.
- Verify functional correctness by confirming both implementations produce identical simulation states.
- Identify and explain specific techniques used to achieve high performance on the GPU.

4. Experimental setup

4.1 Platform

Google Colab is used as the main platform for both implementations. GPU implementation runs on NVIDIA Tesla T4 with a compute capability of 7.5, it has a total memory of 15360 MiB. CPU implementation runs on Colab's provisioned host CPU.

4.2 Grid and simulation parameters

- Grid size: 1024×1024 (1,048,576 cells)
- Generations per trial: 100
- Initial state: pseudo-random, `srand(42)`, ~10% cell density, identical seed used in both implementations for a fair comparison
- Rules: standard Conway B3/S23, applied identically in both

4.3 CPU implementation

Implemented through sequential nested loops over the grid of cells. Edges were checked carefully and compiled it with optimization turned on through “`g++ -O2 -std=c++17.`”

The assigned Colab CPU specifications are as follows,

CPU(s)	2
On-line CPU(s) list	0,1
Model name	Intel(R) Xeon(R) CPU @ 2.20GHz
Thread(s) per core	2
Core(s) per socket	1
Socket(s)	1
NUMA node0 CPU(s)	0,1

Table 1: Colab CPU specifications

4.4 GPU implementation

CUDA thread was used for each cell, with the grid divided into 16×16 blocks. The program is compiled with high optimization through “`nvcc -O3`” and sequentially launches 100 kernels, one per generation, with each kernel computing the entire grid in parallel across GPU threads. Instead of copying data back to CPU after each generation, it swaps the GPU buffers using pointers which makes the process faster.

The assigned Colab GPU specifications are as follows,

GPU	Tesla T4
Compute capability	7.5
Total memory	15,360 MiB
Driver version	580.82.07

Table 2: Colab GPU specifications


4.5 Benchmarking methodology

Both implementations perform a one untimed warm-up pass over the full 100-generation simulation before timing begins. This is done to prime memory allocation, CPU caches and for


GPU context/driver. Afterwards, 5 timed repeats are then run, each resetting the grid to the identical initial state before starting. Wall-clock time for each repeat is measured with “std::chrono::high_resolution_clock” bounded on the GPU side by “cudaDeviceSynchronize()” before and after timing to ensure all kernel work has completed. Mean and range (min-max) across the 5 repeats are reported. Frame-saving for GIF visualization is handled as a separate, untimed run through “saveFrames = true” and is not included in the timing runs used for performance figures (set “saveFrame” to false), since the accompanying disk I/O would otherwise distort the comparison.

4.6 Correctness verification

Alive-cell counts extracted from saved frames at matching generations were compared between the CPU and GPU runs. This was found to be identical which confirms both implementations compute the same simulation.



CPU implementation
98th frame



GPU implementation
98th frame

5. CPU implementation

CPU implementation is done through sequential nested loops over the grid with no parallelism. Every cell is visited one at a time, once per generation. Instead of 2D array of arrays, the grid is stored as a flat “std::vector<int> of size WIDTH × HEIGHT that is indexed through “idx(x, y) = y*WIDTH + x”. This ensures the whole grid is in one contiguous block of memory. This is more cache-friendly than nested pointers/arrays. Cache misses are one of the main causes of slowdown in large grid computations. Therefore, this is an important High-Performance Computing consideration for a single-core implementation.

“countAliveNeighbors()” checks the 8-cell Moore neighbourhood around each cell, with boundary checks so cells at the grid edge only count neighbours that actually exist. “nextState()” applies the B3/S23 rule directly to the current cell's state and neighbour count. Each generation, the grid is updated into a second buffer (next) rather than modified in place. Afterwards, “current.swap(next)” exchanges them in constant time. This avoids the classic bug of a cell's update being influenced by another cell already updated in the same generation.

The initial grid is generated once with a fixed seed (srand(42), ~10% live density) and reused identically across the warm-up pass and all five timed repeats. This ensures every repeat starts from the same input. The program is compiled with “g++ -O2 -std=c++17”. The “-O2” flag enables compiler-level optimizations such as loop unrolling and auto-vectorization. This is the main performance technique applied on the CPU side, since there is no explicit parallelism available to a single core.

The CPU implementation is as follows,

```
%%writefile game_of_life.cpp
#include <iostream>
#include <vector>
#include <chrono>
#include <cstdlib>
#include <ctime>
#include <fstream>
#include <sstream>
#include <iomanip>
#include <algorithm>
#include <numeric>

const int WIDTH = 1024;
const int HEIGHT = 1024;
const int NUM_GENERATIONS = 100;
const int NUM_REPEATS = 5; // run the whole simulation this many times for mean/range

inline int idx(int x, int y) { return y * WIDTH + x; }

int countAliveNeighbors(const std::vector<int>& grid, int x, int y) {
    int alive = 0;
    for (int dy = -1; dy <= 1; ++dy) {
        for (int dx = -1; dx <= 1; ++dx) {
```

```

        if (dx == 0 && dy == 0) continue;
        int nx = x + dx, ny = y + dy;
        if (nx >= 0 && nx < WIDTH && ny >= 0 && ny < HEIGHT) {
            alive += grid[idx(nx, ny)];
        }
    }
}
return alive;
}

int nextState(int current, int aliveNeighbors) {
    if (current == 1 && aliveNeighbors < 2) return 0;
    if (current == 1 && (aliveNeighbors == 2 || aliveNeighbors == 3)) return 1;
    if (current == 1 && aliveNeighbors > 3) return 0;
    if (current == 0 && aliveNeighbors == 3) return 1;
    return current;
}

void savePGM(const std::vector<int>& grid, int generation) {
    std::ostringstream filename;
    filename << "frames/frame_" << std::setw(4) << std::setfill('0') << generation << ".pgm";
    std::ofstream file(filename.str(), std::ios::binary);
    file << "P5\n" << WIDTH << " " << HEIGHT << "\n255\n";
    for (int i = 0; i < WIDTH * HEIGHT; ++i) {
        unsigned char pixel = grid[i] ? 255 : 0;
        file.write(reinterpret_cast<char*>(&pixel), 1);
    }
}

int main() {
    std::srand(42); // fixed seed -> use the same seed in your CUDA version for a fair, identical-input
    comparison

    // Build the initial grid once, then re-copy it fresh before every repeat
    // so every repeat starts from the exact same starting condition.
    std::vector<int> initial(WIDTH * HEIGHT);
    for (auto& cell : initial) cell = (rand() % 10 == 0) ? 1 : 0;

    bool saveFrames = false; // set true only for the visualization run, false for the timing run
    int saveEvery = 2;      // save every 2nd generation -> ~50 frames total

    if (saveFrames) system("mkdir -p frames");

    std::vector<double> times;
    times.reserve(NUM_REPEATS);

    // Warm-up: run one full untimed pass first to prime memory allocation
    // and CPU caches, so the timed repeats below reflect steady-state performance.
    {
        std::vector<int> current = initial;
        std::vector<int> next(WIDTH * HEIGHT);

```

```

    for (int gen = 0; gen < NUM_GENERATIONS; ++gen) {
        for (int y = 0; y < HEIGHT; ++y) {
            for (int x = 0; x < WIDTH; ++x) {
                int alive = countAliveNeighbors(current, x, y);
                next[idx(x, y)] = nextState(current[idx(x, y)], alive);
            }
        }
        current.swap(next);
    }
}

for (int rep = 0; rep < NUM_REPEATS; ++rep) {
    std::vector<int> current = initial; // reset to identical starting grid
    std::vector<int> next(WIDTH * HEIGHT);

    auto start = std::chrono::high_resolution_clock::now();

    for (int gen = 0; gen < NUM_GENERATIONS; ++gen) {
        for (int y = 0; y < HEIGHT; ++y) {
            for (int x = 0; x < WIDTH; ++x) {
                int alive = countAliveNeighbors(current, x, y);
                next[idx(x, y)] = nextState(current[idx(x, y)], alive);
            }
        }
        current.swap(next);

        // only save frames during the first repeat, so extra file I/O
        // doesn't get counted in the other repeats' timings
        if (rep == 0 && saveFrames && gen % saveEvery == 0) savePGM(current, gen);
    }

    auto end = std::chrono::high_resolution_clock::now();
    double elapsed_ms = std::chrono::duration<double, std::milli>(end - start).count();
    times.push_back(elapsed_ms);
    std::cout << "Repeat " << (rep + 1) << "/" << NUM_REPEATS << ": " << elapsed_ms << " ms\n";
}

double sum = std::accumulate(times.begin(), times.end(), 0.0);
double mean = sum / times.size();
double minT = *std::min_element(times.begin(), times.end());
double maxT = *std::max_element(times.begin(), times.end());

std::cout << "\nGrid: " << WIDTH << "x" << HEIGHT << "\n";
std::cout << "Generations: " << NUM_GENERATIONS << "\n";
std::cout << "Repeats: " << NUM_REPEATS << "\n";
std::cout << "Mean total time: " << mean << " ms (range " << minT << " - " << maxT << " ms)\n";
std::cout << "Mean avg time/generation: " << (mean / NUM_GENERATIONS) << " ms\n";

return 0;
}

```

6. CUDA implementation

The CUDA version parallelizes the problem by assigning one GPU thread to every cell. Therefore, all ~ 1 million cells for a generation are evaluated simultaneously instead of sequentially. Instead of looping over cells, the loop is replaced by a single kernel launch where the GPU hardware itself dispatches thousands of threads in parallel. This is the core HPC technique used here.

Each thread computes its own (x, y) from its block and thread indices, and the same neighbour-counting and B3/S23 logic as the CPU version is applied per-thread. The rule itself is unchanged, only the execution model differs. That is what makes the correctness comparison between the two versions meaningful. Threads are organized into 16×16 blocks (256 threads per block, a multiple of the 32-thread warp size, which avoids wasting partially-filled warps), with the grid of blocks sized to cover the full 1024×1024 domain. The boundary check discards the handful of out-of-range threads created by grid dimensions that don't divide evenly.

A second key performance decision is that the grid never leaves GPU memory during the timed simulation loop. The initial grid is uploaded to the device once before timing starts. From then on, “d_current” and “d_next” are swapped by pointer after every kernel launch, so successive generations are computed entirely on-device with no host to device transfer in between. A naive implementation that copied data back to the CPU every generation would be dominated by transfer overhead rather than compute. Avoiding that is one of the main reasons the GPU version is dramatically faster than the CPU one. Kernel launches are wrapped in a “CUDA_CHECK” macro that checks the return code of every CUDA API call, since CUDA errors otherwise fail silently. The program is compiled with “nvcc -O3”.

The GPU implementation is as follows,

```
%%writefile game_of_life_cuda.cu
#include <iostream>
#include <vector>
#include <chrono>
#include <cstdlib>
#include <ctime>
#include <fstream>
#include <sstream>
#include <iomanip>
#include <algorithm>
#include <numeric>
#include <cuda_runtime.h>

const int WIDTH = 1024;
const int HEIGHT = 1024;
const int NUM_GENERATIONS = 100;
const int NUM_REPEATS = 5; // run the whole simulation this many times for mean/range
const int NUM_CELLS = WIDTH * HEIGHT;
```

```

#define CUDA_CHECK(call) \
do { \
    cudaError_t err = call; \
    if (err != cudaSuccess) { \
        std::cerr << "CUDA error at " << __FILE__ << ":" << __LINE__ \
        << " - " << cudaGetErrorString(err) << std::endl; \
        exit(1); \
    } \
} while (0)

__global__ void gameOfLifeKernel(const int* current, int* next, int width, int height) {
    int x = blockIdx.x * blockDim.x + threadIdx.x;
    int y = blockIdx.y * blockDim.y + threadIdx.y;

    if (x >= width || y >= height) return;

    int alive = 0;
    for (int dy = -1; dy <= 1; ++dy) {
        for (int dx = -1; dx <= 1; ++dx) {
            if (dx == 0 && dy == 0) continue;
            int nx = x + dx, ny = y + dy;
            if (nx >= 0 && nx < width && ny >= 0 && ny < height) {
                alive += current[ny * width + nx];
            }
        }
    }

    int idx = y * width + x;
    int state = current[idx];
    int newState = state;

    if (state == 1 && alive < 2) newState = 0;
    else if (state == 1 && (alive == 2 || alive == 3)) newState = 1;
    else if (state == 1 && alive > 3) newState = 0;
    else if (state == 0 && alive == 3) newState = 1;

    next[idx] = newState;
}

void savePGM(const std::vector<int>& grid, int generation) {
    std::ostringstream filename;

```

```

    filename << "frames_cuda/frame_" << std::setw(4) << std::setfill('0') << generation <<
    ".pgm";
    std::ofstream file(filename.str(), std::ios::binary);
    file << "P5\n" << WIDTH << " " << HEIGHT << "\n255\n";
    for (int i = 0; i < WIDTH * HEIGHT; ++i) {
        unsigned char pixel = grid[i] ? 255 : 0;
        file.write(reinterpret_cast<char*>(&pixel), 1);
    }
}

int main() {
    std::srand(42); // same seed as CPU version -> identical initial grid

    std::vector<int> h_initial(NUM_CELLS);
    for (auto& cell : h_initial) cell = (rand() % 10 == 0) ? 1 : 0;

    bool saveFrames = false; // set true only for the visualization run, false for the timing run
    int saveEvery = 2;

    if (saveFrames) system("mkdir -p frames_cuda");

    dim3 blockDim(16, 16);
    dim3 gridDim((WIDTH + blockDim.x - 1) / blockDim.x,
        (HEIGHT + blockDim.y - 1) / blockDim.y);

    // Allocate device buffers once and reuse them across repeats.
    int *d_current, *d_next;
    CUDA_CHECK(cudaMalloc(&d_current, NUM_CELLS * sizeof(int)));
    CUDA_CHECK(cudaMalloc(&d_next, NUM_CELLS * sizeof(int)));

    std::vector<int> h_frame(NUM_CELLS); // only used when saving frames on repeat 0
    std::vector<double> times;
    times.reserve(NUM_REPEATS);

    // Warm-up: launch one full untimed pass first so the timed repeats below
    // don't include one-off CUDA context/driver initialization overhead.
    CUDA_CHECK(cudaMemcpy(d_current, h_initial.data(), NUM_CELLS * sizeof(int),
        cudaMemcpyHostToDevice));
    for (int gen = 0; gen < NUM_GENERATIONS; ++gen) {
        gameOfLifeKernel<<<gridDim, blockDim>>>>(d_current, d_next, WIDTH, HEIGHT);
        std::swap(d_current, d_next);
    }
    CUDA_CHECK(cudaDeviceSynchronize());

```

```

for (int rep = 0; rep < NUM_REPEATS; ++rep) {
    // reset device grid to the identical starting condition before every repeat
    CUDA_CHECK(cudaMemcpy(d_current, h_initial.data(), NUM_CELLS * sizeof(int),
        cudaMemcpyHostToDevice));

    CUDA_CHECK(cudaDeviceSynchronize());
    auto start = std::chrono::high_resolution_clock::now();

    for (int gen = 0; gen < NUM_GENERATIONS; ++gen) {
        gameOfLifeKernel<<<gridDim, blockDim>>>(d_current, d_next, WIDTH, HEIGHT);
        std::swap(d_current, d_next);

        // only save frames during the first repeat, so the device->host copy
        // doesn't get counted in the other repeats' timings
        if (rep == 0 && saveFrames && gen % saveEvery == 0) {
            CUDA_CHECK(cudaMemcpy(h_frame.data(), d_current, NUM_CELLS * sizeof(int),
                cudaMemcpyDeviceToHost));
            savePGM(h_frame, gen);
        }
    }

    CUDA_CHECK(cudaDeviceSynchronize());
    auto end = std::chrono::high_resolution_clock::now();
    double elapsed_ms = std::chrono::duration<double, std::milli>(end - start).count();
    times.push_back(elapsed_ms);
    std::cout << "Repeat " << (rep + 1) << "/" << NUM_REPEATS << ": " << elapsed_ms << "
ms\n";
}

cudaFree(d_current);
cudaFree(d_next);

double sum = std::accumulate(times.begin(), times.end(), 0.0);
double mean = sum / times.size();
double minT = *std::min_element(times.begin(), times.end());
double maxT = *std::max_element(times.begin(), times.end());

std::cout << "\nGrid: " << WIDTH << "x" << HEIGHT << "\n";
std::cout << "Generations: " << NUM_GENERATIONS << "\n";
std::cout << "Repeats: " << NUM_REPEATS << "\n";
std::cout << "Mean total time: " << mean << " ms (range " << minT << " - " << maxT << "
ms)\n";

```

```

std::cout << "Mean avg time/generation: " << (mean / NUM_GENERATIONS) << " ms\n";

return 0;
}

```

7. Comparison results

Metric	CPU (2D array)	CUDA (GPU)
Repeat 1	2404.47 ms	6.41658 ms
Repeat 2	2874.96 ms	6.40737 ms
Repeat 3	3389.84 ms	6.40829 ms
Repeat 4	3003.92 ms	6.41657 ms
Repeat 5	2423.61 ms	6.40888 ms
Mean total time	2819.36 ms	6.41154 ms
Range (min–max)	2404.47 – 3389.84 ms	6.40737 – 6.41658 ms
Mean time/generation	28.19 ms	0.0641 ms

Table 3: Measured execution times (5 repeats, after warm-up)

Using the mean total time across all five repeats, the GPU implementation is approximately **440×** faster than the CPU implementation.

$$\text{Speed up} = \frac{\text{CPU execution time}}{\text{GPU execution time}}$$

$$\text{Speed up} = \frac{2819.36}{6.41154}$$

$$\text{Speed up} = 439.732108 \approx 440$$

The CPU repeats show noticeable variation across the five trials. The repeats 1 and 5 sit close together around 2400–2420 ms, while repeats 2, 3, and 4 are elevated, peaking at 3389.84 ms in repeat 3, which is roughly 41% higher than the fastest repeat. Since the warm-up pass already accounts for first-touch memory allocation and cache-priming costs, this spread more likely reflects background contention on Colab's shared, non-dedicated CPU infrastructure (such as competing processes or other VM tenants on the same physical core). It does not reflect any property of the algorithm itself.

The GPU repeats, by contrast, are tightly clustered within a 0.01 ms band (6.40737–6.41658 ms), essentially indistinguishable from one another. This consistency is itself a meaningful result. Once the GPU context is warmed up and the grid resides entirely in device memory, kernel execution time is highly deterministic, since each kernel launch runs in isolation on dedicated streaming multiprocessors without depending on host OS scheduling.

The scale of the speedup reflects the fundamental difference between the two execution models. The CPU evaluates roughly 1.05 million cells per generation sequentially on a single core, aided only by -O2 compiler optimizations. The CUDA implementation evaluates the same 1.05 million

cells simultaneously across thousands of GPU threads, with memory transfer overhead eliminated by keeping the grid entirely on-device between generations. This confirms that mapping each cell to an independent GPU thread yields a substantial, measurable performance improvement over the sequential CPU baseline which is verified here under identical initial conditions and rules, by using a benchmarking methodology designed to reduce the influence of one-off timing effects.

8. Conclusion

This project implemented Conway's Game of Life on a 1024×1024 grid using two approaches. A sequential CPU implementation using 2D arrays and a parallel CUDA implementation mapping each cell to an independent GPU thread. These methods were benchmarked both under identical, controlled conditions to quantify the performance gained through parallelization.

Both implementations were verified to compute identical simulation states from identical initial conditions. This confirmed that the performance comparison reflects a genuine difference in execution efficiency instead of any difference in what was computed. Using a benchmarking methodology of an untimed warm-up pass followed by five timed repeats, the GPU implementation achieved a mean execution time of 6.41154 ms against the CPU's 2819.36 ms for the same 100-generation simulation. This measured a speedup of approximately $440\times$. This result directly satisfies objectives to quantify, in a statistically reliable way, how much faster the problem can be solved by mapping cells to independent GPU threads.

The results also surfaced a secondary finding beyond the raw speedup figure. The CPU timings varied noticeably across repeats (a 41% spread between the fastest and slowest), most plausibly due to contention on Colab's shared CPU infrastructure, while the GPU timings were highly stable (under a 0.01 ms spread) once warmed up. This asymmetry reinforces that the GPU's performance advantage comes not only from raw throughput but also from more predictable, isolated execution. This is a relevant consideration for HPC workloads where consistency matters as much as speed.

Overall, this project demonstrates the core motivation for High-Performance Computing techniques in practice such as a computationally expensive, embarrassingly parallel problem (recomputing every cell of a large grid every generation). This shows it is impractical to speed up further on a single CPU core, but it maps naturally onto GPU hardware, where thousands of independent threads can each handle one cell's update simultaneously. The specific techniques identified as driving this performance gain were thread-per-cell parallelism, warp-aligned block sizing ($16 \times 16 = 256$ threads/block) and eliminating host–device data transfer from the timed simulation loop by keeping the grid entirely on-device between generations.